

# Learning Domain Randomization Distributions for Transfer of Locomotion Policies


Melissa Mozifian<sup>1,2</sup>, Juan Camilo Gamboa Higuera<sup>1</sup>, David Meger<sup>1,2</sup> and Gregory Dudek<sup>1,2</sup>

<sup>1</sup>: School of Computer Science, McGill University, Montreal, Canada


<sup>2</sup>: Mila – Quebec Artificial Intelligence Institute, Montreal, Canada

## Problem

Learning policies with **low-fidelity** parametric simulation. Unknown dynamics of target system



The AQUA robot simulator



The AQUA robot

What's the best that we can do with **little or no data** from the real system on the desired control task?

## Domain Randomization

Let  $p(z)$  be the distribution of simulation parameters (context)  $z$ . Domain randomization [Tobin et al, 2017] [Peng et al, 2018] aims to find the policy  $\pi$  that maximizes performance over different choices for  $z$

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{p(z)} [J(\pi, z)]$$

$$J(\pi, z) = \mathbb{E}_{p(\tau|\pi, z)} [R(\tau)]$$

where  $R(\tau)$  are the cumulative rewards over trajectory  $\tau$ .

Performance on the target system (e.g. real robot) depends on

- How the policy adapts to different  $z$
- The choice of  $p(z)$


## Choosing the training distribution

To ensure task success, we'd like  $p(z)$  to be as diverse as possible

- The trajectories induced by  $p(z)$  should capture some of the target system dynamics
- But our policy models usually have **finite capacity**

Making  $p(z)$  have high variance does not help

- The set of values for  $z$  where the task is solvable may be small
- But we don't know it a priori



Performance of best policy for hopper with varying  $z$

## Our approach


Pick wide prior  $p(z)$ , where uniform DR may fail.

Use parametric training distribution  $p_{\phi}(z)$ , train policy conditioned on simulation parameters  $\pi_{\theta}(a|s, z)$ . Alternate between:

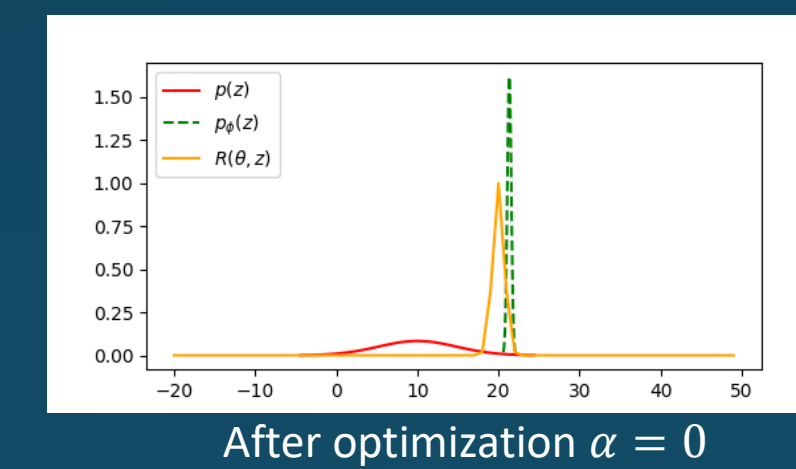
- Optimizing  $\theta$  over the training distribution  $p_{\phi}(z)$
- Optimize  $\phi$  to focus on environments where the current policy performs well
- Use prior to keep from collapsing to easy environments

$$\arg \max_{\theta} \mathbb{E}_{p_{\phi}(z)} [J(\pi_{\theta}, z)]$$


$$\arg \max_{\phi} \mathbb{E}_{p_{\phi}(z)} [J(\pi_{\theta}, z) \log p_{\phi}(z)] - \alpha D_{KL}(p(z) || p_{\phi}(z))$$



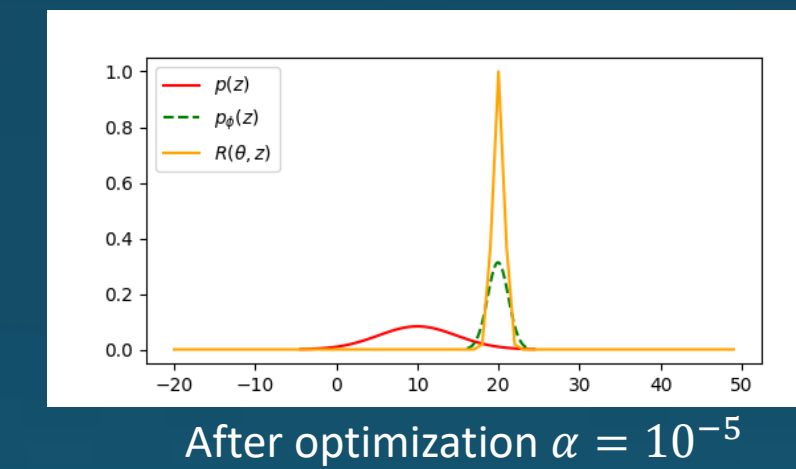
Before optimization



After optimization  $\alpha = 0$



After optimization  $\alpha = 10^{-4}$




After optimization  $\alpha = 10^{-5}$

## Experiments

- Experiments on *Hopper* and *Half-Cheetah* benchmarks [Brockman et al, 2016]
- Modelled  $p_{\phi}(z)$  with discrete distribution,  $p(z)$  with uniform prior
- Compared against uniform DR and using EpOpt-PPO [Rajeswaran, 2016] as policy optimizer

Example results on Hopper:




Learned distributions approximate the ranges found with exhaustive grid evaluations


## Multi-context Distribution Learning

Using rewards seeking objective:

- Multi-context distributions learned for Hopper and Half-Cheetah




Half-Cheetah Learned Context Distribution



Hopper Learned Context Distribution

## Optimizing for worst-case contexts

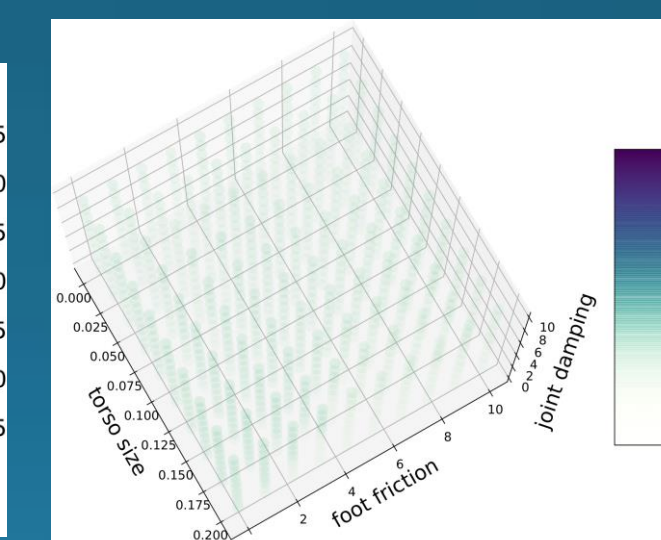


Using the CVaR objective as in EpOpt [Rajeswaran, 2016]

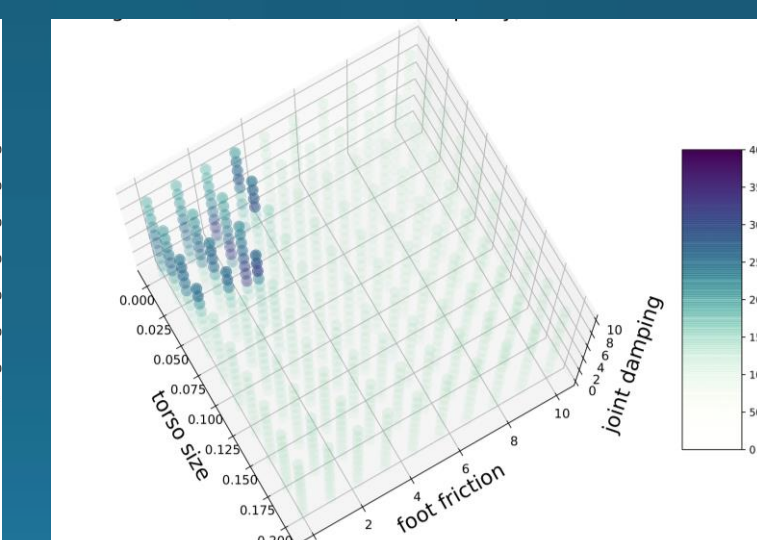
- Learned DR converges faster than uniform DR
- Better asymptotic performance

## Validating the learned distribution


- Performance of policies trained on Hopper over the grid with:



Uniform domain randomization



Learning the parameters of a Gaussian DR distribution



Confidence regions of the learned distribution for 1: 68%, 2: 95%, and 3: 99% confidence